

# CS-E4650 – Homework 3

STEFANO BOSOPPI

ID: 103370763

stefano.bosoppi@aalto.fi

BÁLINT JÁNOS PRÁGAI

ID: 103390468

balintjanos.pragai@aalto.fi

BENEDEK GABOR KISS

ID: 103574543

benedekgabor.kiss@aalto.fi

November 9, 2025

## Contents

<b>1 Methods</b>	<b>2</b>
<b>2 Feature extraction</b>	<b>2</b>
2.1 Numerical features . . . . .	2
2.2 Categorical features . . . . .	3
2.3 Custom migration features . . . . .	3
<b>3 Results</b>	<b>3</b>
3.1 Childcare responsibility division . . . . .	3
3.2 Physical attributes as predictors for lifestyle . . . . .	4
3.3 Biological groups identification . . . . .	4
<b>A Code</b>	<b>5</b>

## 1 Methods

We implemented the data mining pipeline using a shell script. We performed feature extraction from the provided CSV file using a custom program written in Python 3, shown in Appendix A. The `namescodes` utility was used to convert the extracted text-based features into a numerical format for mining and to translate the resulting rules back into a human-readable format.

We conducted association rule discovery, considering both positive and negative ones, using the Kingfisher program. We searched for the top 300 rules. We used Kingfisher's default goodness measure,  $\ln(p_F)$ , as the mutual information measure yielded almost the same results. The other parameter settings included an upper bound for the item number of 170, a lower frequency threshold of 20%, a confidence lower threshold of at least 90%, and an initial goodness measure threshold of -5.

We used the following constraints, as suggested by the task description file:

```

1 eats_insects eats_invertebrates
2 eats_worms eats_invertebrates
3 eats_snails eats_invertebrates
4 eats_larvae eats_invertebrates
5 eats_invertebrates eats_insects eats_worms eats_snails eats_larvae

```

We found no other set of rules that would make sense to exclude and would make a significant difference on the output.

## 2 Feature extraction

We extracted features from the `birds2025ext.csv` file using a Python script, reported in section A, to generate a space-separated transactional file. The extraction logic, detailed below, was designed to convert numerical, categorical, and text-based data into a binary attribute format suitable for association rule mining.

### 2.1 Numerical features

Numerical features were processed to capture the most informative extremes, rather than using the full range of values.

**Intervals** Columns containing ranges (e.g., “10-12”) or single numbers, such as `length`, `wspan`, `weight`, `eggs`, `arrives`, and `leaves`, were first converted to a single floating-point number representing their mean value.

**Derived indices** Two indices, Body Mass Index (BMI) and Wing Span Index (WSI), were calculated from these mean values:

$$BMI = \frac{weight/1000}{(length/100)^2} \quad (1)$$

$$WSI = \frac{wspan}{length} \quad (2)$$

We performed unit conversion for the BMI to respect its definition.

**Discretization** The `mean_eggs`, `BMI`, and `WSI` features were discretized using their first (Q1) and third (Q3) quartiles. Values in the bottom quartile ( $\leq Q1$ ) were assigned a `low_<name>` attribute (e.g., `low_eggs`, `low_wsi`), and values in the top quartile ( $\geq Q3$ ) were assigned a `high_<name>` attribute (e.g., `high_eggs`, `high_bmi`).

## 2.2 Categorical features

Categorical features were converted into binary attributes. The specific logic depended on the feature type, as summarized in Table 1.

Feature type	Original columns	Extraction logic
<b>Taxonomy</b>	group	The group name was added with a group_ prefix (e.g., group_laridae). A hard-coded mapping was used to also add the more general parent group (e.g., group_lari) based on the taxonomy.
<b>Binary (Yes)</b>	diver, long-billed, webbed-feet, long-feet, wading-bird, plunge-dives	The feature name was added to the transaction <i>only</i> if its value was “Yes”. (e.g., webbed_feet).
<b>Binary (Both)</b>	sim	Both “Yes” and “No” values were considered informative. “Yes” was converted to genders_similar and “No” to genders_dissimilar.
<b>Multi-element</b>	diet, biotope	The text was split, and each element was added with a prefix. (e.g., eats_fish, livesin_lakes).
<b>Multi-valued</b>	back, belly, ftype, billcol, legcol, incub, ccare	Each value was split and added with a prefix. (e.g., back_brown, ftype_a, incub_f, ccare_both).

Table 1: Summary of categorical feature extraction logic.

## 2.3 Custom migration features

To capture migration timing, two binary features were created based on the mean arrival and departure months:

- migrates\_arrives\_early: Added if the mean arrival month was  $\leq 4.0$ , meaning April or earlier.
- migrates\_leaves\_late: Added if the mean departure month was  $\geq 10.0$ , meaning October or later.

## 3 Results

From the script output we have identified 4 groups outputs, that are likely to correlate to real-life observations.

### 3.1 Childcare responsibility division

- 1 incub\_both -> genders\_similar fr=31 (0.6200), cf=1.000, gamma=1.250,  
 $\hookrightarrow \Delta=0.124, M=-1.162e+01$
- 2 incub\_both -> ccare\_both fr=28 (0.5600), cf=0.903, gamma=1.328, delta=0.138,  
 $\hookrightarrow M=-1.056e+01$

Based on these two rules we can conclude that if both parents attend the incubation period, they are rather likely to continue with the upbringing of the offspring. In this case the two genders are also similar, likely to camouflage

themselves and the nest while attending the eggs.

- 1 ccare\_f -> incub\_f fr=12 (0.2400), cf=1.000, gamma=2.778, delta=0.154,  
  ↳ M=-1.569e+01
- 2 genders\_dissimilar -> incub\_f fr=10 (0.2000), cf=1.000, gamma=2.778,  
  ↳ delta=0.128, M=-1.237e+01}

Based on the rule “female parent cares for chicks → female parent incubates” is a likely and biologically supported rule, however it would be expected for the reasoning to be the other way around, as the incubating parent would be the likely one to also continue with the upbringing. The other rule is similar in the sense, that it has support in biology, but one would expect reasoning to be in the other direction, however when the two genders are visually distinct it is usually observed that females (who lay the eggs) are the one who incubate them. For this of course they need camouflage colours, while in species like these the males oftentimes have very bright and outstanding colours in order to attract mates.

### 3.2 Physical attributes as predictors for lifestyle

- 1 webbed\_feet -> ~wading\_bird fr=21 (0.4200), cf=0.955, gamma=1.646, delta=0.165,  
  ↳ M=-1.378e+01
- 2 webbed\_feet -> ~eats\_small-rodents fr=22 (0.4400), cf=1.000, gamma=1.250,  
  ↳ delta=0.088, M=-6.663e+00

The first rule highlights a clear functional divergence. Webbed feet are a specific adaptation for propulsion in the water. This is functionally incompatible with the "wading" lifestyle, which typically requires long, un-webbed toes to spread weight on soft mud or move through shallow water with minimal drag. The rule confirms that these two adaptations are mutually exclusive.

The second rule shows that the webbed feet rules out a terrestrial predatory niche. A foot structure optimized for swimming is morphologically unsuited for grasping and killing of land-based prey like rodents. The 100% confidence shows that this ecological boundary is, within this dataset, a strict and unbreakable rule.

### 3.3 Biological groups identification

For this experiment, we lowered the frequency threshold to 10% to have a higher species granularity in the hopes of catching rules defining biological groups.

In the following paragraphs, we present the rules that in our opinion best described three biological groups.

#### Accipitridae

This group, diurnal birds of prey, is defined by a very strong set of interconnected features:

- 1 group\_accipitridae -> ftype\_a fr=8 (0.1600), cf=1.000, gamma=5.556, M=-1.790e+01
- 2 group\_accipitridae -> legcol\_yellow fr=8 (0.1600), cf=1.000, gamma=4.545,  
  ↳ M=-1.500e+01
- 3 ccare\_both legcol\_yellow -> group\_accipitridae fr=8 (0.1600), cf=1.000,  
  ↳ gamma=6.250, M=-2.010e+01

They show that in this dataset, every bird in the group\_accipitridae has flying type A, flapping and soaring flight, and yellow legs. The third rule is even more interesting, showing that the combination of shared parental care and yellow legs is a perfect predictor for this group.

## Anatinae

This group, ducks, is also defined by a strong set of rules, particularly related to appearance and breeding.

```

1 genders_dissimilar high_eggs -> group_anatinae fr=7 (0.1400), cf=1.000,
   ↵ gamma=7.143, M=-1.842e+01
2 webbed_feet ccare_f -> group_anatinae fr=7 (0.1400), cf=1.000, gamma=7.143,
   ↵ M=-1.842e+01
3 group_anatinae -> genders_dissimilar fr=7 (0.1400), cf=1.000, gamma=5.000,
   ↵ M=-1.363e+01
4 group_anatinae -> ccare_f fr=7 (0.1400), cf=1.000, gamma=4.167, M=-1.174e+01

```

These rules have extremely high confidence and lift, linking the biological group to observable differences between the genders, high egg count, and female-only chick care.

It is also interesting to see how a sub-group is distinguished:

```

1 back_dappled genders_dissimilar high_eggs -> group_dabbling_ducks fr=5
   ↵ (0.1000), cf=1.000, gamma=10.000, M=-1.457e+01

```

By adding the `back_dappled` feature to the existing `Anatinae` rules, like `genders_dissimilar` and `high_eggs`, this rule separates dabbling ducks from other `Anatinae`, getting a perfect predictor for this subgroup. The extremely high lift, 10.0, confirms this is a very strong, specific association.

## Scolopacidae

This group, waders, is clearly defined by its feeding habits and related physical traits.

```

1 eats_invertebrates long_billed -> group_scolopacidae fr=5 (0.1000),
   ↵ cf=1.000, gamma=6.250, M=-1.054e+01
2 group_scolopacidae -> eats_invertebrates fr=8 (0.1600), cf=1.000,
   ↵ gamma=2.778, M=-9.415e+00
3 belly_white eats_invertebrates wading_bird -> group_scolopacidae fr=7
   ↵ (0.1400), cf=0.875, gamma=5.469, M=-1.428e+01

```

These rules strongly link the group to its ecological niche. The first rule is a perfect (`cf=1.000`) predictor: a diet of invertebrates, combined with a long bill, is a defining characteristic. The second rule confirms that all birds in this group eat invertebrates. The third rule adds appearance, white belly, and behavior, wading bird, to create a strong, high-lift predictor for the group.

## A Code

Source Code 1: Python code written for the feature extraction.

```

1 import pandas as pd
2 import numpy as np
3 from pathlib import Path
4
5 # Configuration
6 INPUT_FILE = "./birds2025ext.csv"

```

```

7 OUTPUT_FILE = "birdstrans.txt"
8
9 # Taxonomy mapping for hierarchical grouping
10 TAXONOMY_MAP = {
11     'Laridae': 'Lari', 'Sternidae': 'Lari',
12     'Scolopacidae': 'Charadrii', 'Charadriidae': 'Charadrii',
13     'dabbling ducks': 'Anatinae', 'diving ducks': 'Anatinae',
14 }
15
16 BINARY_COLS = ["diver", "long-billed", "webbed-feet", "long-feet",
17     ↵ "wading-bird", "plunge-dives"]
17 MULTI_VALUE_COLS = ["back", "belly", "ftype", "billcol", "legcol", "incub",
18     ↵ "ccare"]
19
20
20 def parse_interval(val):
21     """Convert interval string 'X-Y' to mean, or return float value."""
22     if pd.isna(val):
23         return np.nan
24     s = str(val)
25     if "--" in s:
26         try:
27             parts = s.split("--")
28             return (float(parts[0]) + float(parts[1])) / 2
29         except (ValueError, IndexError):
30             pass
31     try:
32         return float(s)
33     except ValueError:
34         return np.nan
35
36
37 def extract_prefixed_features(text, prefix):
38     """Split multi-element field and return prefixed feature list."""
39     if pd.isna(text):
40         return []
41     elements = str(text).replace(", ", " ").split()
42     return [f"{prefix}_{e.lower().replace(' ', '_')}" for e in elements if e]
43
44
45 def discretize_to_extremes(df, col, feature_name):
46     """Return discretized features for extreme values (Q1/Q3)."""
47     q1, q3 = df[col].quantile([0.25, 0.75])
48
49     def categorize(val):
50         if pd.isna(val):
51             return None
52         if val <= q1:

```

```

53         return f"low_{feature_name}"
54     elif val >= q3:
55         return f"high_{feature_name}"
56     return None
57
58     return df[col].apply(categorize)
59
60
61 def extract_row_features(row, discretized_features):
62     """Extract all features for a single bird species."""
63     features = []
64
65     # Add discretized numerical features
66     for feat in discretized_features:
67         if pd.notna(row[feat]):
68             features.append(row[feat])
69
70     # Taxonomy/Group
71     group = row.get("group")
72     if pd.notna(group):
73         group_clean = str(group).lower().replace(' ', '_')
74         features.append(f"group_{group_clean}")
75     if group in TAXONOMY_MAP:
76         features.append(f"group_{TAXONOMY_MAP[group].lower()}")
77
78     # Diet and Biotope
79     features.extend(extract_prefixed_features(row.get("diet"), "eats"))
80     features.extend(extract_prefixed_features(row.get("biotope"), "livesin"))
81
82     # Binary features (only Yes values)
83     for col in BINARY_COLS:
84         if row.get(col) == "Yes":
85             features.append(col.replace("-", "_"))
86
87     # Gender similarity (both Yes and No)
88     sim = row.get("sim")
89     if sim == "Yes":
90         features.append("genders_similar")
91     elif sim == "No":
92         features.append("genders_dissimilar")
93
94     # Multi-valued categorical features
95     for col in MULTI_VALUE_COLS:
96         val = row.get(col)
97         if pd.notna(val):
98             parts = str(val).lower().replace('-', ' ').split()
99             features.extend(f"{col}_{p}" for p in parts)
100

```

```

101     # Migration timing
102     arrives = row.get("mean_arrives")
103     leaves = row.get("mean_leaves")
104
105     if pd.notna(arrives) and arrives <= 4:
106         features.append("migrates_arrives_early")
107     if pd.notna(leaves) and leaves >= 10:
108         features.append("migrates_leaves_late")
109
110    return sorted(set(features))
111
112
113 def main():
114     # Load data
115     if not Path(INPUT_FILE).exists():
116         print(f"Error: Could not find {INPUT_FILE}")
117         return
118
119     df = pd.read_csv(INPUT_FILE, sep=";", index_col="species")
120
121     # Process interval columns vectorized
122     interval_cols = ["length", "wspan", "weight", "eggs", "arrives", "leaves"]
123     for col in interval_cols:
124         if col in df.columns:
125             df[f"mean_{col}"] = df[col].apply(parse_interval)
126
127     # Calculate derived indices
128     df["BMI"] = (df["mean_weight"] / 1000) / (df["mean_length"] / 100) ** 2
129     df["WSI"] = df["mean_wspan"] / df["mean_length"]
130
131     # Pre-compute discretized features (vectorized)
132     discretized_cols = []
133     for col, name in [("mean_eggs", "eggs"), ("BMI", "bmi"), ("WSI", "wsi")]:
134         if col in df.columns:
135             feat_col = f"disc_{name}"
136             df[feat_col] = discretize_to_extremes(df, col, name)
137             discretized_cols.append(feat_col)
138
139     # Extract features for each row
140     transactions = [extract_row_features(row, discretized_cols) for _, row in
141                     df.iterrows()]
142
143     # Write output
144     with open(OUTPUT_FILE, 'w') as f:
145         for features in transactions:
146             f.write(" ".join(features) + "\n")
147
148     print(f" Feature extraction complete. Saved to {OUTPUT_FILE}")

```

```
148     print(f"  Processed {len(transactions)} species with {sum(len(t) for t in
149         transactions)} total features")
150
151 if __name__ == "__main__":
152     main()
```